


Empirical Security & Privacy, for Humans

UPenn CIS 7000-010



Introduction

Michael Hicks

How would you answer this question?

In the last decade, has the **security** of computer systems, generally,

- improved,
- declined, or
- stayed the same?

A story of memory (un)safety

What is memory safety? - The Programming Languages Enthusiast

www.pl-enthusiast.net/2014/07/21/memory-safety/

The Programming Languages Enthusiast

HOME ABOUT THE PL ENTHUSIAST



← Program verification in the undergraduate CS curriculum

BY MICHAEL HICKS | JULY 21, 2014 · 7:09 AM

Spotlight: Ravi Chugh →

Jump to Comments

What is memory safety?

I am in the process of putting together a [MOOC on software security](#), which goes live in October. At the moment I'm finishing up material on [buffer overflows](#), [format string attacks](#), and other sorts of vulnerabilities in C. After presenting this material, I plan to step back and say, "What do these errors have in common? They are violations of *memory safety*." Then I'll state the definition of memory safety, say why these vulnerabilities are violations of memory safety, and conversely say why memory safety, e.g., as ensured by languages like Java, prevents them.

Recent Posts


- [How to Write a Grad School Personal Statement](#)
- [BullFrog: Online Schema Migration, On Demand](#)
- [Increasing the Impact of PL Research](#)
- ["What is PL Research?" The Talk](#)
- [How to Write a Conference Talk](#)

Subscribe to Blog via Email

Enter your email address to subscribe to this blog and receive notifications of new posts by email.

Violations of memory safety

- Spatial
 - Buffer overflow (heap or stack, read or write)
- Temporal
 - Use after free
 - Use of uninitialized memory
- Other (maybe)
 - Wild pointer deference (int to pointer, deref)
 - Type confusion (bad cast, deref)



Code Red SQL Slammer

Low-level Software Security



History of buffer overflows



The harm has been substantial



- **Morris worm**
 - Propagated across machines (too aggressively, thanks to a bug)
 - One way it propagated was a **buffer overflow** attack against a vulnerable version of `fingerd` on VAXes
 - Sent a special string to the finger daemon, which caused it to execute code that created a new worm copy
 - Didn't check OS: caused Suns running BSD to crash
 - End result: \$10-100M in damages, probation, community service



▶ ▶ 🔍 3:07 / 7:35

Scroll for details

▶ CC ⚙ #

<https://www.youtube.com/watch?v=3BqiTEwz1I0>

What to do? Some options

1. Write C/C++ code without (or with fewer of) these bugs in it
2. Write code in a *memory safe* language
3. Leverage compiler and OS-level *mitigations* that
 1. make the bugs more difficult to exploit, and/or
 2. limit the damage/scope of an exploitation

CERT Secure Coding in C and C++ Professional Certificate

Home > Credentials > CERT Secure Coding in C and C++ Professional Certificate

CERT Secure Coding in C and C++ Professional Certificate


CERT Secure Coding in C and C++ Professional Certificate

The need for qualified experts to support organizations that develop secure software is now greater than ever. To meet this growing demand, we share solutions that are developed as part of our important research. The most effective way to improve software security is to eliminate vulnerabilities during development—before the software is released to users. We offer two certificates in secure coding: Secure Coding in C and C++, described here, and [Secure Coding in Java](#). Both certificates can be earned entirely through online training.

Build More Secure Software

1. Write C/C++ code without (or with fewer of) these bugs in it

The CERT Secure Coding in C and C++ Professional Certificate helps you reduce vulnerabilities in the C and C++ programs they



Register for Certificate Now

Benefits

Who Should Get This Credential?

Term and Renewal

Summary of Fees

How to Earn the

Secure Coding in C and C++

SECOND EDITION



SEI SERIES • A CERT® BOOK

Robert C. Seacord

Foreword by Richard D. Pethia
CERT Director

FREE SAMPLE CHAPTER

scan.coverity.com

BLACKDUCK

Scan Home FAQ OSS Success Stories Projects Using Scan About Community

[Sign up](#) [Sign in](#)

COVERITY SCAN STATIC ANALYSIS

Find and fix defects in your Java, C/C++, C#, JavaScript, Ruby, or Python open source project for free

- ✓ Test every line of code and potential execution path.
- ✓ The root cause of each defect is clearly explained, making it easy to fix bugs
- ✓ Integrated with  

[Sign Up For Free](#)

75% 

CodeQL

codeql.github.com

GitHub Docs Repository License Security Lab

CodeQL

Discover vulnerabilities across a codebase with CodeQL, our industry-leading semantic code analysis engine. CodeQL lets you query code as though it were data. Write a query to find all variants of a vulnerability, eradicating it forever. Then share your query to help others do the same.

CodeQL is free for research and open source.

```
UnsafeDeserialization.ql
import TaintTracking::Global<UnsafeDeserializationConfig>
from PathNode source, PathNode sink
where flowPath(source, sink)
select sink.getNode().(UnsafeDeserializationSink).getMethodAccess(), source, sink,
"Unsafe deserialization of $@.", source.getNode(), "user input"
```

1. Write C/C++ code without (or with fewer of) these bugs in it

GitHub - google/oss-fuzz: OS

github.com/google/oss-fuzz

google / oss-fuzz

Code Issues Pull requests Discussions Actions

master Go to file Code

serge-sans-paille zlib: i... 41ce4da · 6 minutes ago

.allstar Opt out of allstar ... 3 years ago

.clusterfuzzlite ClusterFuzzLite: f... last year

.github Add indexer_build... last month

docs build(deps-dev): ... 4 months ago

infra helper: add ability... 31 minutes ago

projects zlib: improve cove... 6 minutes ago

tools/vscode-exte... build(deps): bum... 4 months ago

.dockerignore [ClusterFuzzLite] ... 4 years ago

.gitattributes Add .gitattributes ... 3 years ago

.gitignore

.pylintrc

About

OSS-Fuzz - continuous fuzzing for open source software.

google.github.io/oss-fuzz

security fuzzing fuzz-testing vulnerabilities stability oss-fuzz

Readme Apache-2.0 license Code of conduct Contributing Security policy Cite this repository Activity Custom properties

1. Write C/C++ code without (or with) OSS-Fuzz annotations

Upstream project

Builder (Cloud Build)

GCS bucket

ClusterFuzz

Developer

Sheriffbot

Issue tracker (monorail)

1. Write fuzzers
2. Commit build configs
3. Sync and build from upstream project
4. Upload to GCS bucket
5. Download and fuzz
6. File bugs, Verify fixes
7. Notify Sheriffbot
8. Fix bugs

Overview

Documentation

Read our [detailed documentation](#) to learn how to use OSS-Fuzz.

Trophies

As of May 2025, OSS-Fuzz has helped identify and fix over 13,000 vulnerabilities and 50,000 bugs across 1,000 projects.

How's that going?

QUANTIFYING MEMORY UNSAFETY AND REACTIONS TO IT

Wednesday, February 03, 2021 - 9:20 am-9:50 am

- **Chrome:** 70% of high/critical vulnerabilities are memory unsafety
- **Firefox:** 72% of vulnerabilities in 2019 are memory unsafety
- **0days:** 81% of in the wild 0days (P0 dataset) are memory unsafe
- **Microsoft:** 70% of all MSRC tracked vulnerabilities are memory unsafety
- **Ubuntu:** 65% of kernel CVEs in USNs in a 6-month sample are memory unsafety
- **Android:** More than 65% of high/critical vulnerabilities are memory unsafety
- **macOS:** 71.5% of Mojave CVEs are due to memory unsafety



an engineer at Mozilla and the United States Digital Service. Alex has a long history of contribution in open source, from building a JIT'd Ruby VM to serving on the Board of Directors of the Python Software Foundation. Alex lives in Washington, D.C.



How's that going?

2024 CWE Top 25					
Rank	ID	Name	Score	CVEs in KEV	Rank Change vs. 2023
1	CWE-79	Improper Neutralization of Input During Web Page Generation ('Cross-site Scripting')	56.92	3	+1
2	CWE-787	Out-of-bounds Write	45.20	18	-1
3	CWE-89	Improper Neutralization of Special Elements used in an SQL Command ('SQL Injection')	35.88	4	0
4	CWE-352	Cross-Site Request Forgery (CSRF)	19.57	0	+5
5	CWE-22	Improper Limitation of a Pathname to a Restricted Directory ('Path Traversal')	12.74	4	+3
6	CWE-125	Out-of-bounds Read	11.42	3	+1
7	CWE-78	Improper Neutralization of Special Elements used in an OS Command ('OS Command Injection')	11.30	5	-2
8	CWE-416	Use After Free	10.19	5	-4
9	CWE-862	Missing Authorization	10.11	0	+2
10	CWE-434	Unrestricted Upload of File with Dangerous Type	10.03	0	0
11	CWE-94	Improper Control of Generation of Code ('Code Injection')	7.13	7	+12

Classic memory-safety vulnerabilities

12	CWE-20	Improper Input Validation	6.78	1	-6
13	CWE-77	Improper Neutralization of Special Elements used in a Command ('Command Injection')	6.74	4	+3
14	CWE-287	Improper Authentication	5.94	4	-1
15	CWE-269	Improper Privilege Management	5.22	0	+7
16	CWE-502	Deserialization of Untrusted Data	5.07	5	-1
17	CWE-200	Exposure of Sensitive Information to an Unauthorized Actor	5.07	0	+13
18	CWE-863	Incorrect Authorization	4.05	2	+6
19	CWE-918	Server-Side Request Forgery (SSRF)	4.05	2	0
20	CWE-119	Improper Restriction of Operations within the Bounds of a Memory Buffer	3.69	2	-3
21	CWE-476	NULL Pointer Dereference	3.58	0	-9
22	CWE-798	Use of Hard-coded Credentials	3.46	2	-4
23	CWE-190	Integer Overflow or Wraparound	3.37	3	-9
24	CWE-400	Uncontrolled Resource Consumption	3.23	0	+13
25	CWE-306	Missing Authentication for Critical Function	2.73	5	-5

What to do? Some options

1. Write C/C++ code without (or with fewer of) these bugs in it
2. Write code in a *memory safe* language
3. Leverage compiler and OS-level *mitigations* that
 1. make the bugs more difficult to exploit, and/or
 2. limit the damage/scope of an exploitation

2. Write code in a *memory safe* language

Memory safe:

- Rust
- Swift
- Go
- Haskell
- Python
- Etc.

Memory unsafe:

- C
- C++
- Assembly

How's that going?

the percentage of memory safety vulnerabilities in Android dropped from 76% to 24% over 6 years as development shifted to memory safe languages

The screenshot shows a web browser window displaying the Google Security Blog. The title of the article is "Eliminating Memory Safety Vulnerabilities at the Source", posted on September 25, 2024. The post discusses the transition to memory-safe languages in Android, stating that the percentage of memory safety vulnerabilities dropped from 76% to 24% over 6 years. It highlights the use of Safe Coding as a secure-by-design approach. The browser interface includes standard navigation buttons, a search bar, and a tab for the Google Online Security Blog.

Google Online Security Blog

security.googleblog.com/2024/09/eliminating-...

Google Security Blog

The latest news and insights from Google on security and safety on the Internet

Eliminating Memory Safety Vulnerabilities at the Source

September 25, 2024


Posted by Jeff Vander Stoep - Android team, and Alex Rebert - Security Foundations

Memory safety vulnerabilities remain a pervasive threat to software security. At Google, we believe the path to eliminating this class of vulnerabilities at scale and building high-assurance software lies in [Safe Coding](#), a secure-by-design approach that prioritizes transitioning to memory-safe languages.



This post demonstrates why focusing on Safe Coding for new code quickly and counterintuitively reduces the overall security risk of a codebase, finally breaking through the stubbornly high plateau of memory safety vulnerabilities and starting an exponential decline, all while being scalable and cost-effective.

We'll also share updated data on how the percentage of memory safety vulnerabilities in Android dropped from 76% to 24% over 6 years as development shifted to memory safe languages.

How's that going?



How's that going?





Australian Government
Australian Signals Directorate

ASD
AUSTRALIAN SIGNALS DIRECTORATE
Australian Cyber Security

Communications Security Establishment
Canadian Centre for Cyber Security

Centre de la sécurité des télécommunications
Centre canadien pour la cybersécurité

National Cyber Security Centre
a part of GCHQ

National Cyber Security Centre
PART OF THE GCHQ



The Case for Memory Safe Roadmaps

Why Both C-Suite Executives and Technical Experts Need to Take Memory Safe Coding Seriously

Publication: December 2023

United States Cybersecurity and Infrastructure Security Agency

United States National Security Agency

United States Federal Bureau of Investigation

Australian Signals Directorate's Australian Cyber Security Centre

Canadian Centre for Cyber Security

United Kingdom National Cyber Security Centre

New Zealand National Cyber Security Centre

Computer Emergency Response Team New Zealand

December 2023

This document is marked TLP:CLEAR. Disclosure is not limited. Sources may use TLP:CLEAR when information carries minimal or no foreseeable risk of misuse, in accordance with applicable rules and procedures for public release. Subject to standard copyright rules, TLP:CLEAR information may be distributed without restriction. For more information on the Traffic Light Protocol, see cisa.gov/tlp.

TLP:CLEAR



Australian Government
Australian Signals Directorate

ASD
AUSTRALIAN SIGNALS DIRECTORATE
Australian Cyber Security

Communications Security Establishment
Canadian Centre for Cyber Security

Centre de la sécurité des télécommunications
Centre canadien pour la cybersécurité



Exploring Memory Safety in Critical Open Source Projects

Publication: June 26, 2024

Cybersecurity and Infrastructure Security Agency (CISA)

Federal Bureau of Investigation (FBI)

Australian Signals Directorate's (ASD's) Australian Cyber Security Centre (ACSC)

Canadian Centre for Cyber Security (CCCS)

June 2024

This document is marked TLP:CLEAR. Disclosure is not limited. Sources may use TLP:CLEAR when information carries minimal or no foreseeable risk of misuse, in accordance with applicable rules and procedures for public release. Subject to standard copyright rules, TLP:CLEAR information may be distributed without restriction. For more information on the Traffic Light Protocol, see cisa.gov/tlp.

What to do? Some options

1. Write C/C++ code without (or with fewer of) these bugs in it
2. Write code in a *memory safe* language
3. Leverage compiler and OS-level *mitigations* that
 1. make the bugs more difficult to exploit, and/or
 2. limit the damage/scope of an exploitation

3. Leverage compiler and OS-level *mitigations* that
 1. make the bugs more difficult to exploit, and/or
 2. limit the damage/scope of an exploitation

Challenge exploitability:

- Stack canaries
- Address-space layout randomization (ASLR)
- “write xor execute” ($W \oplus X$)
- Control-flow integrity (CFI)
- Etc.

Limit damage:

- Process-level isolation
- Within-process compartments (eg., RLbox)
- Externally enforced access control

How's that going?

A screenshot of a web browser window. The title bar says "Project Zero: The More You K...". The address bar shows "googleprojectzero.blogspot.com/2022/04/the-mor...". The main content area has a large heading "Project Zero" and a sub-section "News and updates from the Project Zero team at Google".

Project Zero

News and updates from the Project Zero team at Google

Tuesday, April 19, 2022

The More You Know, The More You Know You Don't Know

A Year in Review of 0-days Used In-the-Wild in 2021

Posted by Maddie Stone, Google Project Zero

This is our third annual year in review of 0-days exploited in-the-wild [[2020](#), [2019](#)]. Each year we've looked back at all of the detected and disclosed in-the-wild 0-days as a group and synthesized what we think the trends and takeaways are. The goal of this report is not to detail each individual exploit, but instead to analyze the exploits from the year as a group, looking for trends, gaps, lessons learned, successes, etc. If you're interested in the analysis of individual exploits, please check out our [root cause analysis repository](#).

We perform and share this analysis in order to **make 0-day hard**. We want it to be more costly, more

A screenshot of a web browser window. The title bar says "Project Zero: Oday "In the Wild" x". The address bar shows "googleprojectzero.blogspot.com/p/0day.html". The main content area has a large heading "Project Zero" and a sub-section "News and updates from the Project Zero team at Google". Below this is a specific post titled "0day "In the Wild"".

0day "In the Wild"

Posted by Ben Hawkes, Project Zero (2019-05-15)

Project Zero's team mission is to "make zero-day hard", i.e. to make it more costly to discover and exploit security vulnerabilities. We primarily achieve this by performing our own security research, but at times we also study external instances of zero-day exploits that were discovered "in the wild". These cases provide an interesting glimpse into real-world attacker behavior and capabilities, in a way that nicely augments the insights we gain from our own research.

Today, we're sharing our tracking spreadsheet for publicly known cases of detected zero-day exploits, in the hope that this can be a useful community resource:


Spreadsheet link: [0day "In the Wild"](#)

This data is collected from a range of public sources. We include relevant links to third-party analysis and attribution, but we do this only for your information; their inclusion does not mean we endorse or validate the content there. The data described in the spreadsheet is nothing new, but we think that collecting it together in one place is useful. For example, it shows that:

- On average, a new "in the wild" exploit is discovered every 17 days (but in practice these often clump together in exploit chains that are all discovered on the same date);
- Across all vendors, it takes 15 days on average to patch a vulnerability that is being used in active attacks;
- A detailed technical analysis on the root-cause of the vulnerability is published for 86% of listed CVEs;
- Memory corruption issues are the root-cause of 68% of listed CVEs.

We also think that this data poses an interesting question: *what is the detection rate of 0day exploits?* In other words, at what rate are 0day exploits being used in attacks *without* being detected? This is a key "unknown parameter" in security, and how you model it will greatly inform your views, plans, and priorities as a defender.

How's that going?



How's that going?

Project Zero: The More You K
x


← → C googleprojectzero.blogspot.com/2022/04/the-mor... ☆ 🔒 📁 🔍

New Year, Old Techniques

We had a record number of “data points” in 2021 to understand how attackers are actually using 0-day exploits. A bit surprising to us though, out of all those data points, there was nothing new amongst all the data. 0-day exploits are considered one of the most advanced attack methods an actor can use, so it would be easy to conclude that attackers must be using special tricks and attack surfaces. But instead, the 0-days we saw in 2021 generally followed the same bug patterns, attack surfaces, and exploit “shapes” previously seen in public research. Once “0-day is hard”, we’d expect that to be successful, attackers would have to find new bug classes of vulnerabilities in new attack surfaces using never before seen exploitation methods. In general, that wasn’t what the data showed us this year. With two exceptions (described below in the iOS section) out of the 58, everything we saw was pretty “[meh](#)” or standard.

Out of the 58 in-the-wild 0-days for the year, 39, or 67% were memory corruption vulnerabilities. Memory corruption vulnerabilities have been the standard for attacking software for the last few decades and it’s still how attackers are having success. Out of these memory corruption vulnerabilities, the majority also stuck with very popular and well-known bug classes:

- 17 use-after-free
- 6 out-of-bounds read & write
- 4 buffer overflow
- 4 integer overflow



Out of the 58 in-the-wild 0-days for the year, 39, or 67% were memory corruption vulnerabilities.

How's that going?

Lots we don't know

- No estimate of the *damage* caused by the exploitation
- Data is *lower bound* – probably more exploited zero-days than these (not known, not reported)
- Zero-days are *not the only vector of attack*

2025 Data Breach Investigation Report

verizon.com/business/resources/reports/dbir/

verizon business


Log In

Contact sales

Call Sales: 844-669-0847

2025 Data Breach Investigations Report


Today's threat landscape is shifting. Get the latest updates on real-world breaches and help safeguard your organization from cybersecurity attacks.



Key resources Testimonial Top takeaways Webinars Sign up Archive FAQs

Key resources

Feedback



2025 DBIR

2025 DBIR Executive Summary

2025 DBIR infographic





Figure 1. Known initial access vectors in non-Error, non-Misuse breaches (n=9,891)

The exploitation of vulnerabilities has seen another year of growth as an initial access vector for breaches, reaching 20%. This value approaches that of credential abuse, which is still the most common vector. This was an increase of 34% in relation to last year's report and was supported, in part, by zero-day exploits targeting edge devices and virtual private networks (VPNs). The percentage of edge devices and VPNs as a target on our exploitation of vulnerabilities action was 22%, and it grew almost eight-fold from the 3% found in last year's report. Organizations worked very hard to patch those edge device vulnerabilities, but our analysis showed only about 54% of those were fully remediated throughout the year, and it took a median of 32 days to accomplish.

The screenshot shows the homepage of the report. At the top, there's a navigation bar with links for 'Log In', 'Contact sales', and 'Call Sales: 844-669-0847'. Below the header, the title '2025 Data Breach Investigations Report' is prominently displayed in red. A sub-headline reads 'Today's threat landscape is shifting. Get the latest updates on real-world breaches and help safeguard your organization from cybersecurity attacks.' To the right, there's a large, stylized red and yellow graphic resembling a shield or a 3D cube. Below the title, there's a section titled 'Key resources' with three thumbnail images: '2025 DBIR', '2025 DBIR Executive Summary', and '2025 DBIR infographic'.

The exploitation of vulnerabilities has seen another year of growth as an initial access vector for breaches, reaching 20%. This value approaches that of credential abuse, which is still the most common vector.

The US Treasury Department was hacked



Hugo Herrera / The Verge

The Treasury Department said a China-based threat actor gained access to several employee workstations and unclassified documents.

by Emma Roth

Dec 30, 2024, 5:25 PM EST

25 Comments (25 New)

beyondtrust.com/remote-support-saaS-service-security-investigation?utm_source=google&utm_medium=cpc&utm_campaign=remote-support-saas-service-security-investigation&utm_content=remote-support-saas-service-security-investigation&utm_term=remote-support-saas-service-security-investigation

Security Incident Details

BeyondTrust confirmed and began taking measures to address the security incident on December 5, 2024 that involved our Remote Support SaaS product. No BeyondTrust products outside of Remote Support SaaS were affected. No FedRAMP instances were affected. No other BeyondTrust systems were compromised, and ransomware was not involved.

Our investigation into the cause and impact of the compromise was conducted with a recognized third-party cybersecurity and forensics firm. The investigation determined that a zero-day vulnerability of a third-party application was used to gain access to an online asset in a BeyondTrust AWS account. Access to that asset then allowed the threat actor to obtain an infrastructure API key that could then be leveraged against a separate AWS account which operated Remote Support infrastructure. This vulnerability, as well as the two vulnerabilities discovered and disclosed as noted in the timeline above have been patched.

... “gain access to an online asset in a BeyondTrust AWS account. Access to that asset then allowed the threat actor to **obtain an infrastructure API key** ...” which was used to operate the Remote Support infrastructure

AT&T says criminals stole phone records of ‘nearly all’ customers in new data breach

Breach linked to Snowflake

Zack Whittaker

Snowflake blamed the data thefts on its customers for not using multi-factor authentication to secure their Snowflake accounts, a security feature that the cloud data giant did not enforce or require its customers to use.

Cybersecurity incident response firm Mandiant, which Snowflake called in to help with notifying customers, later said about 165 Snowflake customers had a “significant volume of data” stolen from their customer accounts.

“Snowflake blamed the data thefts on its customers for **not using multi-factor authentication** to secure their snowflake accounts, ... **did not require its customers to use”**

The XZ Backdoor: Everything You Need to Know

Details are starting to emerge about a stunning supply chain attack that sent the open source software community reeling.




ILLUSTRATION: DA-KUK/GETTY IMAGES

“This might be the best executed **supply chain attack** we've seen described in the open, and it's a nightmare scenario: **malicious, competent, authorized upstream in a widely used library**”

The big picture



<https://www.usenix.org/conference/usenixsecurity19/presentation/stamos>




28TH USENIX
SECURITY SYMPOSIUM

Open Access Sponsor



The big picture



<https://www.usenix.org/conference/usenixsecurity19/presentation/stamos>



28TH USENIX
SECURITY SYMPOSIUM

Open Access Sponsor




So: What would you do?

- If you were a CISO, or a VP of Security Engineering, how would you spend your money?
- If you were the head of a government agency like CISA, tasked with improving the state of cybersecurity, what would you recommend?
- If you were the head of NSF's Secure and Trustworthy Cyberspace program, what would you fund?

Empirical Security & Privacy, for Humans

UPenn CIS 7000-010



Goals for the course

- Learn the research state of the art

- Economic view of cybersecurity
 - End users and cybersecurity
 - Cybersecurity as a scientific pursuit
 - Cybersecurity investment as risk assessment
 - Cybersecurity game theory
 - Cyberattack economics
 - Cybersecurity public health
-
- Developers' and operators' actions, and security
 - Ethics in computer security experimentation
 - Network-based security measurement
 - Privacy

Goals for the course

- Learn the research state of the art
- Learn relevant research methods
- Learn how to learn – how to dive into a field, understand its results, and see gaps and opportunities
- Do something interesting: New result, reproduction, or a deep dive

<https://canvas.upenn.edu/courses/1880676>

The screenshot shows a web browser window for the course CIS 7000-010 202530 Empirical Security & Privacy, For Humans on the canvas.upenn.edu platform. The URL in the address bar is canvas.upenn.edu/courses/1880676. The course title is displayed prominently on the right. A sidebar on the left contains links for Account, Dashboard, Courses (which is selected), Calendar, Inbox, History, and Help. The main content area includes a search bar, course details (BAN_CIS-7000-010 202530), a navigation menu with Home, Assignments, Discussions, Grades, People, Pages, Syllabus, BigBlueButton, Collaborations, and Search, and a detailed course description.

CIS 7000-010 202530 Empirical Security & Privacy, For Humans

This graduate seminar course has two main goals.

1. To understand ways to evaluate security meaningfully. Security is a tradeoff, imposing a cost to enable a benefit, and it is important to assess this tradeoff carefully. A theorem or an in-lab demonstration is necessary but not sufficient.
2. To understand how innovators should think about technical security and privacy solutions in realistic usage, i.e., used not by people like the designer, but rather “normal” people in practical settings.

This is a graduate seminar, focused on reading technical literature and discussing and exploring ideas.

See the [syllabus](#) for details on course content, and grading.

Approach

- Read papers, (sometimes) present them, critique them, discuss
 - Learn from experts in the field
 - Do a (group) project
-
- We will have Zoom enabled during the class, for remote lectures and for those who (occasionally) can't make it

Read papers: Question, understand, improve

How to Read a Paper

S. Keshav
David R. Cheriton School of Computer Science, University of Waterloo
Waterloo, ON, Canada
keshav@uwaterloo.ca

ABSTRACT

Researchers spend a great deal of time reading research papers. However, this skill is rarely taught, leading to much wasted effort. This article outlines a practical and efficient *three-pass method* for reading research papers. I also describe how to use this method to do a literature survey.

Categories and Subject Descriptors: A.1 [Introductory and Survey]

General Terms: Documentation.

Keywords: Paper, Reading, Hints.

1. INTRODUCTION

Researchers must read papers for several reasons: to review them for a conference or a class, to keep current in their field, or for a literature survey of a new field. A typical researcher will likely spend hundreds of hours every year reading papers.

Learning to efficiently read a paper is a critical but rarely taught skill. Beginning graduate students, therefore, must learn on their own using trial and error. Students waste much effort in the process and are frequently driven to frustration.

4. Glance over the references, mentally ticking off the ones you've already read

At the end of the first pass, you should be able to answer the *five Cs*:

1. *Category*: What type of paper is this? A measurement paper? An analysis of an existing system? A description of a research prototype?
2. *Context*: Which other papers is it related to? Which theoretical bases were used to analyze the problem?
3. *Correctness*: Do the assumptions appear to be valid?
4. *Contributions*: What are the paper's main contributions?
5. *Clarity*: Is the paper well written?

Using this information, you may choose not to read further. This could be because the paper doesn't interest you, or you don't know enough about the area to understand the paper, or that the authors make invalid assumptions. The *first pass* is adequate for papers that aren't in your research

Class prep

- Read the paper(s) for that class. Submit a 1-2 paragraph review

Reading assignment, Aug 28 | +

← → ⌂ canvas.upenn.edu/courses/1880676/assignments/13860302

Search this course

The screenshot shows a Canvas course page for 'BAN_CIS-7000-010 ...' with the title 'Reading assignment, Aug 28 (part I)'. The assignment is due Wednesday by 3pm, worth 5 points, and allows submissions via text entry or file upload. It specifies doc, txt, and pdf file types and is available after Aug 25 at 10:30am. The page includes a sidebar with various course links like Home, Assignments, Discussions, Grades, People, Pages, Syllabus, BigBlueButton, Collaborations, Search, and Help.

202530 (Fall 2025)

Home

Assignments

Discussions

Grades

People

Pages

Syllabus

BigBlueButton

Collaborations

Search

File Upload Text Entry

Copy and paste or type your submission right here.


Edit View Insert Format Tools Table

12pt ▾ Paragraph ▾ :

Guest lectures (so far)



Alex Gantman
VP, Security Engineering,
Qualcomm, August 28



Cormac Herley
VP, Security Engineering,
Qualcomm, September 4



Adam Shostack
Founder & CEO, Shostack
& Associates, October 7

Present papers: Distill, reveal, dive deep


- Will do this for the second half of the class
- We will vote on a pool of papers to present, and you can select the 1 you want
- Grading criteria: Understanding, thoughtfulness, background/perspective, clarity, materials quality, delivery, non-regurgitation, answering questions

Projects

- Something substantial: New study, reproduction, literature review, ...
- Timeline
 - Pitches in class @ 9/25
 - Proposal @ 10/9
 - Final paper @ finals week

About me

- Ph.D., CIS @ UPenn 2001
- Remained a Philly sports fan (go Eagles!)



About me

Employment

- 2002-2022 – TTk faculty, UMD
- 2006-2015 – Adjunct, IDA/CCS (NSA-funded research lab)
- 2008, 2015 – Visiting Researcher, Microsoft Research
- 2018-2021 – CTO, Correct Computation Inc (startup)
- 2022-present – Scientist, Scientific Computing and Numerical Algorithms Group, University of Maryland



COMPUTER SCIENCE
UNIVERSITY OF MARYLAND



Leadership positions

- 2002-2022 – Member, DARPA Information Science and Technology Board



ector, Maryland Center


- 2002-2022 – Chair, ACM Special Group on Programming Languages



Education, UMD CS

About me

- Research @ UMD: **Software Security**, Programming Languages, Software Engineering, Usability, Cryptography, Quantum Computing, Networks, Databases
- Startup: Building tools for secure software development
 - Binary analysis
 - Migration to memory-safe C
- AWS
 - Cedar authorization language
 - Fuzzing/automated test generation
 - Formal/mechanized proofs of security



Cedar: a new authorization language

Focuses on **centralized** decision-making



Powers **Amazon Verified Permissions** and **AWS Verified Access**
Powers **StrongDM** and **Common Fate** access solutions

Open source at
<https://github.com/cedar-policy>

Reading for next time

Plus: “How to Read a Paper?”

Why Information Security is Hard – An Economic Perspective

Ross Anderson

University of Cambridge Computer Laboratory,
JJ Thomson Avenue, Cambridge CB3 0FD, UK
Ross.Anderson@ccl.cam.ac.uk

Abstract

According to one common view, information security comes down to technical measures. Given better access control policy models, formal proofs of cryptographic protocols, approved firewalls, better ways of detecting intrusions and malicious code, and better tools for system evaluation and assurance, the problems can be solved.

In this note, I put forward a contrary view: information insecurity is at least as much due to perverse incentives. Many of the problems can be explained more clearly and convincingly using the language of microeconomics: network externalities, asymmetric information, moral hazard, adverse selection, liability dumping and the tragedy of the commons.


1 Introduction

In a survey of fraud against autoteller machines [4], it was found that patterns of fraud depended on who was liable for them. In the USA, if a customer disputed a transaction, the onus was on the bank to prove that the customer was mistaken or lying; this gave US banks a motive to protect their systems properly. But in Britain, Norway and the Netherlands, the burden of proof lay on the customer; the bank was right up

risk of forged signatures from the bank that relies on the signature (and that built the system) to the person alleged to have made the signature. Common Criteria evaluations are not made by the relying party, as Orange Book evaluations were, but by a commercial facility paid by the vendor. In general, where the party who is in a position to protect a system is not the party who would suffer the results of security failure, then problems may be expected.

A different kind of incentive failure surfaced in early 2000, with distributed denial of service attacks against a number of high-profile web sites. These exploit a number of subverted machines to launch a large coordinated packet flood at a target. Since many of them flood the victim at the same time, the traffic is more than the target can cope with, and because it comes from many different sources, it can be very difficult to stop [7]. Varian pointed out that this was also a case of incentive failure [20]. While individual computer users might be happy to spend \$100 on anti-virus software to protect themselves against attack, they are unlikely to spend even \$1 on software to prevent their machines being used to attack Amazon or Microsoft.

This is an example of what economists refer to as the ‘Tragedy of the Commons’ [15]. If a hundred peasants on their land in the hillside共同耕种一片草地，



The screenshot shows a web browser window titled "The Market for Silver Bullets". The URL in the address bar is "iangu.org/papers/market_for_silver_bullets.html". The page content includes the title "The Market for Silver Bullets" and author "Ian Grigg Systems, Inc.". It also shows the date "2nd March 2008". Below the title, there is a section titled "Abstract: What is security?" followed by the text of the paper. The text discusses the concept of security as a "good" in economics, noting that it is often a safety good that is destructive and hard to test. It also mentions the "silver bullet" hypothesis and its relation to the "Job Market Signaling" theory.

The Market for Silver Bullets

Ian Grigg
Systems, Inc.

2nd March 2008

Abstract: What is security?

As a “good” in the sense of economics, security is now recognised as being one for which our knowledge is poor. As with safety goods, events of utility tend to be destructive, yet unlike safety goods, the performance of the good is very hard to test. The roles of participants are complicated by the inclusion of aggressive attackers, and buyers and sellers that interchange.

This essay hypothesises that security is a good with insufficient information, and rejects the assumption that security fits in the market for goods with asymmetric information. Security can be viewed as a market where neither buyer nor seller has sufficient information to be able to make a rational buying decision. Drawing heavily from Michael Spence’s “Job Market Signaling,” these characteristics lead to the arises of a market in *silver bullets* as participants *herd* in search of *best practices*, a common set of goods that arises more to reduce the costs of externalities rather than achieve benefits in security itself.

Introduction

In an investigation into security, Adam Shostack posed the question, *what are good signals in the market for security* [1] [2]? In addressing this apparently clear question we find ourselves drawn to the question of *what is security*? One avenue of potential investigation is to ask what the science of economics can provide in answer to this question. In economics terms, security could be a “good” as it is demanded and traded for value. This essay seeks to cast security as a good, and attempts to classify what sort of good it is?